

Programming with Python for Engineers

Release 0.0.1

Sinan Kalkan, Onur Tolga Sehitoglu, Gokturk Ucoluk

Oct 07, 2020

Contents

Preface	1
Basic Computer Organization	9

Preface

i. About this book

a. Target audience of the book

This book is intended to be an accompanying textbook for teaching programming to science and engineering students with no prior programming expertise. This endeavour requires a delicate balance between providing details on computers & programming in a complete manner and the programming needs of science and engineering disciplines. With the hopes of providing a suitable balance, the book uses Python as the programming language, since it is easy to learn and program. Moreover, for keeping the balance, the book is formed of three parts:

- **Part I: The Basics of Computers and Computing:** The book starts with what computation is, introduces both the present-day hardware and software infrastructure on which programming is performed and introduces the spectrum of programming languages.
- **Part II: Programming with Python:** The second part starts with the basic building blocks of Python programming and continues with providing the ground formation for solving a problem in to Python. Since almost all science and engineering libraries in Python are written with an object-oriented approach, a gentle introduction to this concept is also provided in this part.
- **Part III: Using Python for Science and Engineering Problems:** The last part of the book is dedicated to practical and powerful tools that are widely used by various science and engineering disciplines. These tools provide functionalities for reading and writing data from/to files, working with data (using e.g. algebraic, numerical or statistical computations) and plotting data. These tools are then utilized in example problems and applications at the end of the book.

b. How to use the book

This is an ‘interactive’ book with a rather ‘minimalist’ approach: Some details or specialized subjects are not emphasized and instead, direct interaction with examples and problems are encouraged. Therefore, rather than being a ‘complete reference manual’, this book is a ‘first things first’ and ‘hands on’ book. The pointers to skipped details will be provided by links in the book. Bearing this in mind, the reader is strongly encouraged to read and interact all contents of the book thoroughly.

The book’s interactivity is thanks to [Jupyter notebook](https://jupyter.org)¹. Therefore, the book differs from a conventional book by providing some dynamic content. This content can appear in audio-visual form as well as some applets (small applications) embedded in the book. It is also possible that the book asks the the reader to complete/write a piece of Python program, run it, and inspect the result, from time to time. The reader is

¹ <https://jupyter.org>

encouraged to complete these minor tasks. Such tasks and interactions are of great assistance in gaining acquaintance with Python and building up a self-confidence in solving problems with Python.

Thanks to Jupyter notebook running solutions on the Internet (e.g. [Google Colab](https://colab.research.google.com/)², [Jupyter Notebook Viewer](https://nbviewer.jupyter.org/)³), there is absolutely no need to install any application on the computer. You can directly download and run the notebook on Colab or Notebook Viewer. Though, since it is faster and it provides better virtual machines, the links to all Jupyter notebooks will be served on Colab.

ii. What is computing?

Computing is the process of inferring data from data. What is going to be inferred is defined as the *task*. The original data is called the *input (data)* and the inferred one is the *output (data)*.

Let us look at some examples:

- Multiplying two numbers, X and Y , and subtracting 1 from the multiplication is a *task*. The two numbers X and Y are the *input* and the result of $X \times Y - 1$ is the *output*
- Recognizing the faces in a digital picture is a *task*. Here the *input* is the color values (3 integers) for each point (pixel) of the picture. The *output* is, as you would expect, the pixel positions that belong to faces. In other words, the output can be a set of numbers.
- The board instance of a chess game, as *input*, where black has made the last move. The task is to predict the best move for white. The best move is the *output*.
- The *input* is a set of three-tuples which look like **<Age_of_death, Height, Gender>**. The *task*, an optimization problem in essence, is to find out the curve (i.e. the function) that goes through these tuples in a 3D dimensional space spanned by Age, Height and Gender. As you have guessed already, the *output* is the parameters defining the function and an error describing how well the curve goes through the tuples.
- The *input* is a sentence to a chatbot. The *task* is to determine the sentence (the *output*) that best follows the input sentence in a conversation.

These examples suggest that computing can involve different types of data, either as input or output: Numbers, images, sets, or sentences. Although this variety might appear intimidating at first, we will see that, by using some ‘solution building blocks’, we can do computations and solve various problems with such a wide spectrum of data.

Certainly not! This is a common mistake a layman does. There are diverse architectures based on totally different physical phenomena that can compute. A good example is the *brain* of living beings, which rely on completely different mechanisms compared to the *micro processors* sitting in our laptops, desktops, mobile phones and calculators.

The building blocks of a *brain* is the *neuron*, a cell that has several input channels, called *dendrites* and a single output channel, the *axon*, which can branch like a tree (see [Fig. 1](#)).

² <https://colab.research.google.com/>

³ <https://nbviewer.jupyter.org/>




Fig. 1: Our brains are composed of simple processing units, called neurons. Neurons receive signals (information) from other neurons, process those signals and produce an output signal to other neurons. [Drawing by BruceBlaus - Own work, CC BY 3.0, <https://commons.wikimedia.org/w/index.php?curid=28761830>]

The branches of an axon, each carrying the same information, connect to other neurons' dendrites (Fig. 2). The connection with another neuron is called the *synapse*. What travels through the synapse are called *neurotransmitters*. Without going into details, one can simplify the action of neurotransmitters as messengers that cause an excitation or inhibition on the receiving end. In other words, the neurotransmitters, through a chemical process along the axon, are released into the synapse as the 'output' of the neuron, they 'interact' with the dendrite (i.e. the 'input') of another neuron and potentially lead to an excitation or an inhibition.




Fig. 2: Neurons ‘communicate’ with each other by transmitting neurotransmitters via synapses. [Drawing by user:Looie496 created file, US National Institutes of Health, National Institute on Aging created original - <http://www.nia.nih.gov/alzheimers/publication/alzheimers-disease-unraveling-mystery/preface>, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=8882110>]

Interestingly enough, the synapse is like a valve, which reduces the neurotransmitters’ flow. We will come

to this in a second. Now, all the neurotransmitters flow in through the input channels (dendrites) have an accumulative effect on the (receiving) neuron. The neuron emits a neurotransmitter burst through its axon. This emission is not a ‘what-comes-in-goes-out’ type. It is more like the curve in Fig. 3.




Fig. 3: When a neuron receives ‘sufficient’ amount of signals, i.e. stimulated, it emits neurotransmitters on its axon, i.e. it fires. [Plot by Original by en:User:Chris 73, updated by en:User:Diberri, converted to SVG by tiZom - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=2241513>]

The throughput of the synapse is something that may vary with time. Most synapses have the ability to ease the flow over time if the neurotransmitter amount that entered the synapse was constantly high. High activity widens the synaptic connection. The reverse also happens: Less activity over time narrows the synaptic connection.

Some neurons are specialized in creating neurotransmitter emission under certain physical effects. Retina neurons, for example, create neurotransmitters if light falls on them. Some, on the other hand, create physical effects, like creating an electric potential that will activate a muscle cell. These specialized neurons feed the huge neural net, the brain, with inputs and receive outputs from it.

The human brain, containing about 10^{11} such neurons with each neuron being connected to 1000-5000 other neurons by the mechanism explained above, is a very unique computing ‘machine’ that inspires computational sciences.

A short video on synaptic communication between neurons

The brain never stops processing information and the functioning of each neuron is only based on signals (the neurotransmitters) it receives through its connections (dendrites). There is no common synchronization timing device for the computation: i.e. each neuron behaves on its own and functions in parallel.

An interesting phenomenon of the brain is that the information and the processing are distributed. Thanks to this feature, when a couple of neurons die (which actually happens each day) no information is lost com-

pletely.

On the contrary to the brain, which uses varying amounts of chemicals (neurotransmitters), the microprocessor based computational machinery uses the existence and absence of an electric potential. The information is stored very locally. The microprocessor consists of subunits but they are extremely specialized in function and far less in number compared to 10^{11} all alike neurons. In the brain, changes take place at a pace of 50 Hz maximum, whereas this pace is 10^9 Hz in a microprocessor.

In Chapter 1, we will take a closer look at the microprocessor machinery which is used by today's computers. Just to make a note, there are man-made computing architectures other than the microprocessor. A few to mention would be the 'analog computer', the 'quantum computer' and the 'connection machine'.

iv. What is a 'computer'?

As you have already noticed, the word 'computer' is used in more than one context.

1. **The broader context:** Any physical entity that can do 'computation'.
2. **The most common context:** An electronic device that has a 'microprocessor' in it.

From now on, 'computer' will refer to the second meaning, namely a device that has a 'microprocessor'.

A computer...

- is based on binary (0/1) representations such that all inputs are converted to 0s and 1s and all outputs are converted from 0/1 representations to a desired form, mostly a human-readable one. The processing takes places on 0s and 1s, where 0 has the meaning of 'no electric potential' (no voltage, no signal) and 1 has the meaning of 'some fixed electric potential (usually 5 Volts, a signal).
- consists of two clearly distinct entities: The Central Processing Unit (CPU), also known as the microprocessor (\square P), and a *Memory*. In addition to these, the computer is connected to or incorporates other electronic units, mainly for input-output, known as 'peripherals'.
- performs a 'task' by executing a sequence of instructions, called a 'program'.
- is deterministic. That means if a 'task' is performed under the same conditions, it will produce always the same result. It is possible to include randomization in this process only by making use of a peripheral that provides electronically random inputs.

v. What is programming?

The CPU (the microprocessor - \square P) is able to perform several types of actions:

- Arithmetic operations on binary numbers that represent (encode) integers or decimal numbers with fractional part.
- Operations on binary representations (like shifting of digits to the left or right; inverting 0s and 1s).
- Transferring to/from memory.
- Comparing numbers (e.g. whether a number n_1 larger than n_2) and performing alternative actions based on such comparisons.
- Communicating with the peripherals.

- Alternating the course of the actions.

Each such unit action is recognized by the CPU as an *instruction*. In more technically terms, tasks are solved by a CPU by executing a sequence of instructions. Such sequences of instructions are called machine codes. Constructing machine codes for a CPU is called ‘machine code programming’.

But, programming has a broader meaning:

a series of steps to be carried out or goals to be accomplished.

And, as far as computer programming is concerned, we would certainly like these steps to be expressed in a more natural (more human readable) manner, compared to binary machine codes. Thankfully, there exist ‘machine code programs’ that read-in such ‘more natural’ programs and convert them into ‘machine code programs’ or immediately carry out those ‘naturally expressed’ steps.

Python is such a ‘more natural way’ of expressing programming steps.

Basic Computer Organization

In this chapter, we will provide an overview of the internals of a modern computer. To do so, we will first describe a general architecture on which modern computers are based. Then, we will study the main components and the principles that allow such machines to function as general purpose “calculators”.

The von Neumann Architecture

Basic Computer Organization. In this chapter, we will provide an overview of the internals of a modern computer. To do so, we will first describe a general architecture on which modern computers are based. Then, we will study the main components and the principles that allow such machines to function as general purpose “calculators”.

John von Neumann

From: [Oxford Reference](#)⁴

“Hungarian-born US mathematician, creator of the theory of games and pioneer in the development of the modern computer. Born in Budapest, the son of a wealthy banker, von Neumann was educated at the universities of Berlin, Zürich, and Budapest, where he obtained his PhD in 1926. After teaching briefly at the universities of Berlin and Hamburg, von Neumann moved to the USA in 1930 to a chair in mathematical physics at Princeton. In 1933, he joined the newly formed Institute of Advanced Studies at Princeton as one of its youngest professors. By this time he had already established a formidable reputation as one of the most powerful and creative mathematicians of his day. In 1925 he had offered alternative foundations for set theory, while in his *Mathematischen Grundlagen der Quantenmechanik* (1931) he removed many of the basic doubts that had been raised against the coherence and consistency of quantum theory. In 1944, in collaboration with Oskar Morgenstern (1902–77), von Neumann published *The Theory of Games and Economic Behaviour*. A work of great originality, it is reputed to have had its origins at the poker tables of Princeton and Harvard. The basic problem was to show whether it was possible to speak of rational behaviour in situations of conflict and uncertainty as in, for example, a game of poker or wage negotiations. In 1927 von Neumann proved the important theorem that even in games that are not fully determined, safe and rational strategies exist. With entry of the USA into World War II in 1941 von Neumann, who had become an American citizen in 1937, joined the Manhattan project (for the manufacture of the atom bomb) as a consultant. In 1943 he became involved at Los Alamos on the crucial problem of how to detonate an atom bomb. Because of the enormous quantity of computations involved, von Neumann was forced to seek mechanical aid. Although the computers he had in mind could not be made in 1945, von Neumann and his colleagues began to design Maniac I (Mathematical analyser, numerical integrator, and computer). Von Neumann was one of the first to see the value of a flexible stored program: a program that could be changed quite easily without altering the computer’s basic circuits. He went on to consider deeper problems in the theory of logical automata and

⁴ <https://www.oxfordreference.com/view/10.1093/oi/authority.20110803120234729>

finally managed to show that self-reproducing machines were theoretically possible. Such a machine would need 200 000 cells and 29 distinct states. Having once been caught up in affairs of state von Neumann found it difficult to return to a purely academic life. Thereafter much of his time was therefore spent, to the regret of his colleagues, advising a large number of governmental and private institutions. In 1954 he was appointed to the Atomic Energy Commission. Shortly after this, cancer was diagnosed and he was forced to struggle to complete his last work, the posthumously published *The Computer and the Brain* (1958).”



Fig. 4: John von Neumann (1903 – 1957)

Components of the von Neumann Architecture

The von Neumann architecture (Fig. 5) defines the basic structure, or outline, used in most computers today. Proposed in 1945 by Von Neumann, it consists of two distinct units: An *addressable memory* and a *Central Processing Unit* (CPU). All the encoded actions and data are stored together in the memory unit. The CPU, querying these actions, the so-called *instructions*, executes them one by one, sequentially (though, certain instructions may alter the course of execution order).

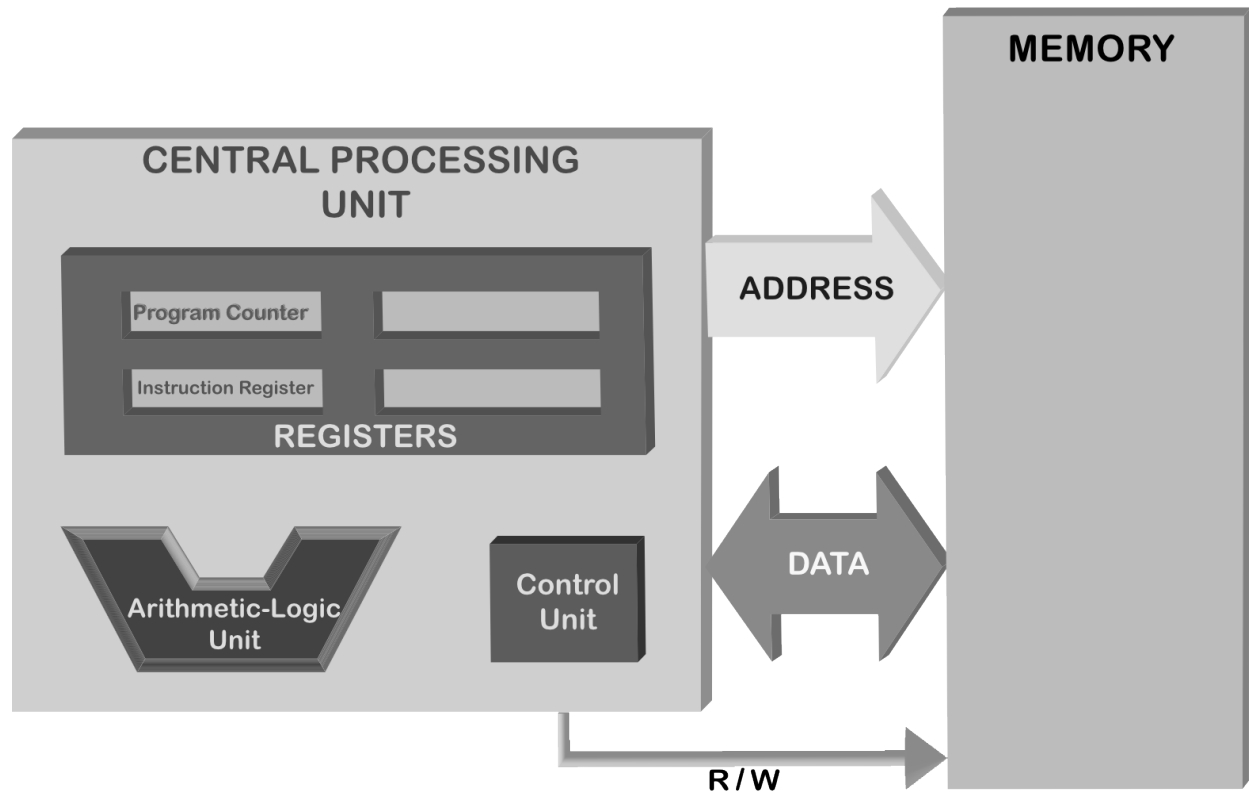


Fig. 5: A block structure view of the von Neumann Architecture.

The CPU communicates with the memory by two sets of wires, namely the *address bus* and the *data bus*, plus a single *R/W* wire (Fig. 5). These busses consist of several wires and carry a binary information to/from the memory. Each wire in a bus carries one bit of the information (either a zero (0) or a one(0)). Today's von Neumann architectures are working on electricity, therefore, these zeros and ones correspond to Voltages. A one is usually the presence of a 5V and a zero is the absence of it.

The Memory

The memory can be imagined as pigeon holes organized as rows (Fig. 6). Each row has eight pigeon holes, each being able to hold a zero (0) or one (1) – in electronic terms, each of them is capable of storing a Voltage (can you guess what type of an electronic component a pigeon hole is?). Each such row is named to be of the size *byte*. In computer terms, a byte means 8 bits.




Fig. 6: The memory is organized as a stack of rows such that each row has an associated address.

Each byte of the memory has a unique address. When the address input (also called address bus – Fig. 5) of the memory is provided a binary number, the memory byte that has this number as the address becomes accessible through the data output (also called output data bus). Based on W/R wire being set to Write (1) or Read (0), the action that is carried out on the memory byte differs:

- **W/R wire is set to WRITE (1) :**

The binary content on the input data bus is copied into the 8-bit location whose address is provided on the address bus, the former content is *overwritten*.

- **W/R wire is set to READ (0) :**

The data bus is set to a copy of the content of 8-bit location whose address is provided on the address bus. The content of the accessed byte is left intact.

The information stored in this way at several addresses live in the memory happily, until the power is turned off.

The memory is also referred as Random Access Memory (RAM). Some important aspects of this type of memory has to be noted: * Access any content in RAM, whether for reading or writing purposes, is *only* possible when the content's address is provided to the RAM through the address bus. * Accessing any content takes exactly the same amount of time, irrespective of the address of the content. In todays RAMs, this access time is around 50 nanoseconds. * When a content is overwritten, it is gone forever and it is not possible to undo this action.

An important question is who sets the address bus and who communicates through the data bus (sends and receives bytes of data). As depicted in Fig. 5, the CPU does. How this is done on the CPU side will become

clear in the section below.

The CPU

The Central Processing Unit, which can be considered as the ‘brain’ of a computer, consists of the following units:

- **Control Unit (CU)**, which is responsible for fetching instructions from the memory, interpreting (‘decoding’) them and executing them. After executing an instruction finishes, the control unit continues with the next instruction in the memory. This “fetch-decode-execute” cycle is constantly executed by the control unit.
- **Arithmetic Logic Unit (ALU)**, which is responsible for performing arithmetic (addition, subtraction, multiplication, division) and logic (less-than, greater-than, equal-to etc.) operations. CU provides the necessary data to ALU and the type of operation that needs to be performed, and ALU executes the operation.
- **Registers**, which are mainly storage units on the CPU for storing the instruction being executed, the affected data, the outputs and temporary values.

The size and the quantity of the registers differ from CPU model to model. They generally have size in a range of [2-64] bytes and today’s most popular CPUs most registers have size 64 bits (i.e. 8 bytes). Their quantity is not high and in the range of [10-20]. The registers are broadly categorized into two: *Special Purpose Registers* and *General Purpose Registers*.

Two special purpose registers are worth mentioning to understand how a CPU’s Fetch-Decode-Execute cycle runs. The first is the so-called *Program Counter* (PC) and the second is the *Instruction Register* (IR).

- **Inputs/Outputs connections**, which connect the CPU to the other components in the computer.

The Fetch-Decode-Execute Cycle

The CPU is in fact a *state machine*, a machine that has a representation of its current *state*, reads the next instruction and executes the instruction according to its current state. The state consists what is stored in the registers. Until it is powered off, it follows the Fetch-Decode-Execute cycle (Fig. 7) where each step of the cycle is based on its state. The *control unit* is responsible for delicate cycle.




Fig. 7: The CPU constantly follows the fetch-decode-execute cycle while the computer is running a program.

- **The Fetch Phase**

The cycle starts with the Fetch Phase. The CPU has the address (the position in the memory) of the next instruction in the PC (Program Counter) Register. During this phase, the address bus is set to this value in the PC register and the R/W wire is set to Read (0). The memory responds to this by providing the memory content at the given address on the data bus.

How many bytes are sent through the data bus? That is architecture dependent. Usually it is 4-8 bytes. These bytes are received into the IR (Instruction Register).

- **The Decode Phase**

The content of the first part of IR electronically triggers some action. Every CPU has an electronically built-in hard-wired instruction table in which every possible atomic instruction that the CPU can carry out has an associated binary code, called *operation code* (opcode in short). This table differs from CPU brand to brand.

An example for such an atomic instruction (simply called as *instruction*) is integer addition on two registers. Prior to the instruction, the two registers contain integers, and after the instruction is executed, one of the registers will be incremented by the amount of the other (by means of integer addition).

In the decode phase, the electronic activation of the specific electronic circuitry, a subpart of the ALU, that will carry out the instruction is performed. In contrary to the ‘addition’ example we just introduced, some instructions may require additional data. For example, loading a register with a certain integer or floating point value; or transferring a result obtained in a register to a certain place in the memory. Where does this additional data come from? (the integer or floating point in the first case; the address of the place that will receive the resulting content of the register, in the second case). This type of data adjunct to an instruction follows right after the instruction byte.

This is illustrated below in an 8-byte instruction. In this example, the designer of the CPU allocated the first four bits for representing the type of operation the instruction is supposed to execute (e.g. the designer could design the CPU such that 0001 means that this is an instruction for reading data from memory) and the remaining four bits contain the affected data (in this case, the address of the memory content we are going to read).

Opcode	Affected data or address
0001	0100

There are three types of instructions:

- *Data manipulation*: Arithmetic/Logic operations on/among registers,
- *Data transfer*: Memory-to-Register, Register-to-Memory, Register-to-Register transfers,
- *Control flow of execution*: Instructions that stop execution, jump to a different part of the memory for next instruction, instead of the next one in the memory.
- **The Execute Phase**

As the name implies, the electronically activated and initialized circuitry carries out the instruction. Depending on the instruction, the registers, the memory or other components are affected. When the instruction completes, the PC is updated by one unless it was a control flow changing instruction in which case PC is updated to point to the to-be-jumped address in the memory. Not all instructions take the same amount of time to be carried out. Floating point division, for example, takes much more time compared to others.

A CPU goes through the *Fetch-Decode-Execute* cycle until it is powered off. What happens at the very beginning? The electronics of the CPU is manufactured such that, when powered up, the PC register has a very fixed content. Therefore, the first instruction is always fetched from that position.

An intelligent question would be “when does the CPU jump from one state to another?”. One possible answer is: whenever the previous state is completed electronically, a transition to the next state is performed. Interestingly, this is *not* true. The reality is that there is an external input to the CPU from which electronic pulses are fed. This input is called the *system clock* and each period of it is named as a *clock cycle*. The best performance would be that each phase of the fetch-decode-execute cycle is completed in one-and-only-one clock cycle. On modern CPUs, this is true for addition instruction, for example. But there are instructions (like floating point division) which take about 40 clock cycles.

What is length of a clock cycle? CPUs are marked with their *clock frequency*. For example, Intel’s latest processor, i9, has a maximal clock frequency of 5GHz (that is 5×10^9 pulses per second). So, since (period) = $1/(\text{frequency})$, for this processor a clock cycle is 200 pico seconds. This is such a short time that light would travel only 6 cm.

A modern CPU has many more features and functional components: *interrupts*, *ports*, *various levels of caches* are a few of them. To cover them is certainly out of the scope of this course material.

The Stored Program Concept

In order for the CPU compute something, the corresponding instructions to do the computation have to be placed into the memory (how this is achieved will become clear in the next chapter). These instructions and data that perform a certain task are called a *Computer Program*. The idea of storing a computer program into the memory to be executed is coined as the *Stored Program Concept*.

What does a stored program look like? Below you see a real extract from the memory, a program that multiplies two integer numbers sitting in two different locations in the memory and stores the result in another memory location (to save space consecutive 8 bytes in the memory are displayed in a row, the next row displays the next 8 bytes):


```

01010101 01001000 10001001 11100101 10001011 00010101 10110010 00000011
00100000 00000000 10001011 00000101 10110000 00000011 00100000 00000000
00001111 10101111 11000010 10001001 00000101 10111011 00000011 00100000
00000000 10111000 00000000 00000000 00000000 00000000 11001001 11000011
...
11001000 00000001 00000000 00000000 00000000 00000000

```

Unless you have a magical talent, this should not be understandable to you. It is difficult because it is just a sequence of bytes. Yes the first byte is presumably some instruction, but what is it? Furthermore, since we do not know what it is, we do not know whether it is followed by some data or not, so we cannot say where the second instruction starts. However, the CPU for which these instructions were written for would know this, hard-wired in its electronics.

When a programmer wants to write a program at this level, i.e. in terms of binary CPU instructions and binary data, s/he has to understand and know each instruction the CPU can perform, should be able to convert data to some internal format, make a detailed memory layout on paper and then start writing down each bit of the memory. This is an extremely painful job; though it is possible, it is impractical.

Alternatively, consider the text below:

```

pushq %rbp
    movq %rsp, %rbp
    movl alice(%rip), %edx
    movl bob(%rip), %eax
    imull %edx, %eax
    movl %eax, carol(%rip)
    movl $0, %eax
    leave
    ret
alice:
    .long 123
bob:
    .long 456

```

Though `pushq` and `movq` are not immediately understandable, the rest of the text provides some hints. `alice` and `bob` must be some programmer's name invention, something like variables with values 123 and 456 respectively; `imull` must have something to do with 'multiplication', since only registers can be subject to arithmetic operations `%edx` and `%eax` must be some denotation used for registers; having uncovered this, `movl`s start to make some sense: they are some commands to move around data... and so on. Even without knowing the instruction set, with a small mind gaming we can uncover the action sequence.

This text is an example *assembly* program. A human invented denotation for instructions and data. An important piece of knowledge is that each line of the assembler text corresponds to a single instruction. This assembly text is so clear that even manual conversion to the cryptic binary code above is feasible. From now on, we will call the binary code program as a *Machine Code Program* (or simply the *machine code*).

How do we automatically obtain machine codes from assembly text? We have machine code programs that convert the assembly text into machine code. They are called *Assemblers*.

Despite making programming easier for programmers, compared to machine codes, even assemblers are insufficient for efficient and fast programming. They lack some high-level constructs and tools that are necessary for solving problems easier and more practical. Therefore higher level languages that are much easier to read and write compared to assembly are invented.

We will cover the spectrum of programming languages in more detail in the next chapter.

Pros and Cons of the von Neuman Architecture

Advantages

- CPU retrieves data and instruction in the same manner from a single memory device. This simplifies the design of the CPU.
- Data from input/output (I/O) devices and from memory are retrieved in the same manner. This is achieved by mapping the device communication electronically to some address in the memory.
- The programmer has a considerable control of the memory organization. So, s/he can optimize the memory usage to its full extend.

Disadvantages

- Sequential instruction processing nature makes parallel implementations difficult. Any parallelization is actually a quick sequential change in tasks.
- The famous “*Von Neumann bottleneck*” : Instructions can only be carried out one at a time and sequentially.
- Risk of an instruction being unintentionally overwritten due to an error in the program.

Peripherals of a computer

Though it is somewhat contrary to your expectation any device outside of the Von Neumann structure, namely the CPU and the Memory is a *peripheral*. In this aspect even the keyboard, the mouse and the display are peripherals. So are the USB and ethernet connection and the internal hard disk. To study the technical details of how those devices are connected to the Von Neumann architecture is out of the scope of this book. Though we can summarize it in a few words.

All devices are electronically listening to the busses (the address and data bus) and to a wire running out of the CPU (which is *not* pictured above) is 1 or 0. This wire is called the *port_io* line and tells the memory devices as well as to any other device that listens to the busses whether the CPU is talking to the (real) memory or not. If it is talking to the memory all the other listeners keep extreme silent. But if the *port_io* line is 1, meaning the CPU doesn't want to talk to the memory but to the device which is electronically sensitive to that specific address that was put on the address bus (by the CPU), then that device jumps up and responds (through the data bus). The CPU can send as well as receive data from that particular device. A computer has some precautions to prevent address clashes, i.e. two devices responding to the same address information in *port_io*. Another mechanism aids communication requests initiated from the peripherals. Of course it would be possible for the CPU from time to time stop and do a *port_io* on all possible devices, asking them for any data they want to send in. This technique is called *polling* and is extremely inefficient for devices that send asynchronous data (data that is send in irregular intervals). You cannot know when there will be a keyboard entry so, in polling, you have to ask very frequently the keyboard device for the existence of any data. Instead of the dealing with the inefficiency of polling another mechanism is built into the CPU. The interrupt mechanism is an electronic circuitry of the CPU which has inlets (wires) connected to the peripheral devices. When a device wants to communicate with (send or receive some data to/from) the CPU then send a signal (1) from that specific wire. This gets the attention of the CPU, the CPU stops what it is doing at a convenient point in time, and ask the device for a *port_io*. So the device gets a chance to send/receive data to/from the CPU.

The running of a computer

When you power on a computer, it first goes through a start-up process (also called booting), which, after performing some routine checks, loads a program from your disk called Operating System.

Start up Process

At the core of a computer sits a Von Neumann architecture. But how finds a machine code its way into the memory, gets settled there, so the CPU starts executing it is still unclear. It is obvious that even when you buy a brand new computer and turn it on for the first time it does some actions which is traceable on its display. So there is a machine code in the memory which, even when the power is off, remains. Or, of course, there is a portion of the memory which is supplied with electricity even when the computer is turned off. It is the first option. There is a portion of the memory which does not lose its content, very much like a flash drive. It is electronically located exactly at the address where the CPU looks for its first instruction. This portion of the memory, with its content, is called Basic Input Output System, or in short BIOS. In the former days the BIOS was manufactured as write-only-once. To change the program a chip had to be replaced with a new one. The size of the BIOS of the first PCs was 16KB, nowadays it is about 1000 times larger, 16MB. When you power up a PC the BIOS program will do the following in sequence:

- Power-On Self Test, abbreviated as POST, determines whether the CPU and memory is intact, identifies and if necessary initializes devices like the video display card, keyboard, hard disk drive, optical disc drive and other basic hardware.
- In a predefined order locate boot loader software on “boot devices”. It is possible to set some parameters of the BIOS, one of which is the boot devices (e.g. a hard disk, a floppy disk, a USB flash stick, CD, or DVD) and in which order they will be queried. BIOS will get the Master Boot Record (MBR) from the first available device in the given order. MBR is supposed to contain a short machine code program that will load the *Operating System* (OS). So, BIOS loads the content of the MBR into the memory and start to execute it. This program loads the actual operating system and then starts the execution of it.

The Operating System

The operating system is a program that after being loaded into the memory manages resources and services like the use of memory, CPU and devices. The OS has components that manage these resources which are explained below:

- **Memory Management:** Refers to the management of the Memory connected to the CPU. In modern computers there are more than one machine code programs loaded into the memory. Each device for some task. Some of them are initiated by the user (like a browser, document editor, excel, music player, etc.) and some are initiated by the operating system at the boot up of the computer. The CPU switches very fast from one program in the memory (this is called as *process*) to another. The user (usually) does not feel the switching. The memory manager, keeps track of the space allocated by processes in the memory. When a new program (process) is being started it has to be placed into the memory. The memory manager decides where it is going to be placed. Of course when the process ends the place in the memory occupied by it has to be reclaimed, that is the memory managers job. It is also possible, while a process is running, it demands some additional space in the memory (e.g. a photoshop-like program needs space for a newly opened jpg image) then the process makes this demand to the memory manager, which grants it or denies it.

- **Processor (Time) Management:** As said above, in modern computers more than one machine code programs are loaded into different locations of the memory. An electronic mechanism forces the CPU to switch to the *Time Manager* component of the OS. At least 20 times a second the time manager is invoked and makes a decision on behalf of the CPU. This of the processes that sit in the memory will be run for the next period? When a process gets the turn, the current state of the CPU (content of all its registers) is saved to some secure position in the memory, in association to the last executing process. From that secure position the last saved state information which going to take the turn is found and the CPU is set to that state. Then the CPU, for a period of time executed that process. At the end of that period, CPU switches over to the time manager and the time manager makes a decision for the next period. Which process will get the turn? And so on. This decision making is a complex task. Still there are Ph.D. level research going on on this subject. The time manager takes some statistics about each individual process about its system resource utilization. Also there is the possibility that a process has a high priority associated due to several reasons. The time manager has to solve a kind of optimization problem under some constraints. As mentioned this is a complex task and a hidden quality factor of an OS.
- **Device Management:** All external devices of a computer have a software plug-in to the operating system. An operating system has some standardized demands from devices and these software plug-ins implement these standardized functionality. This software is usually provided by the device manufacturer and is loaded into the operating system as a part of the device installing process. These plug-ins are named as *device drivers*.

An Operating System performs device communication by means of these drivers. It does the following activities for device management:

- Keeps tracks of all devices' status.
 - Decides which process gets access to the device when and for how long.
 - Implements some intelligent caching, if possible, for the data communication with the device.
 - De-allocates devices.
- **File Management:** A computer is basically a data processing machine. Various data are produced or used for very diverse purposes. Textual, numerical, audio-visual data is handled. One kind of handling of data is *storing* and *retrieving* it on some external recording device. Examples to such recording devices are hard disks, flash drives, CDs and DVDs. Data is stored on devices under files. A *file* is a persistent piece of information that has a name, some meta data (data about the data like owner, creation data, size, content type, etc) and the data. The organizational way how files are stored on devices is called the *file system*. There are various alternatives to do this. FAT16, FAT32, NTFS, EXT2, EXT3, ExFAT, HFS+ are a few of about a hundred (actually the most common ones). Each of them have their pros and cons as far as *max allowed file size*, *security*, *robustness (repairability)*, *extensibility*, *metadata*, *lay out policies* and some more are concerned. Files are most often managed in a hierarchy. This is achieved by a concept of directories. On the surface (as far as the user sees them) a file system usually consist of files separated into groups, so called *directories*, where directories can contain files or other directories.

The file manager is responsible of creation, initialization of a file system, inclusion and removal of devices from this system and management of all sort of alternation in the file system: Creation, removal, copying of files and directories, dynamically assigning access rights for files and directories to processes are among them.

- **Security:** This is basically the maintenance of system integrity, availability, and confidentiality. The security of a computer exists at various layers such as
 - maintaining the physical security of the computer system,

- the security of the information the system is in hold of,
- and the security of the network to which the computer is connected.

In all of these areas, the operating system plays a vital role in keeping the security. Especially the second item is where the operating system is involved at most. Information, as you know by now, is placed in the computer in two locations. The internal memory and the external storing devices. The internal memory is in hold of the processes and the processes shall not interfere with each other (unless specifically intended). Actually in a modern day computer there can be more than one user working at the same time on the computer. Their processes running in the memory as well as their files on the file system must remain extremely private. Even their existence have to be hidden from every other user.

Computers are connected and more and more integrated in a global network. This integration is done on a spectrum of interactions. In the extreme case a computer can be solely controlled over the network. Of course this is done by supplying some credentials but, as you would guess, such a system is prone to malicious attacks. An operating system has to take measures to protect the computer system from such harms. Sometimes it is the case that bugs of the OS is discovered and exploited in order to breach security.

- **User Interface:** As the computer's user, when you want to do anything you do this by ordering the operating system to do it. Starting/terminating a program, exploring or modifying the file system, installing/uninstalling a new device are all done by talking to the operating system. For this purpose an OS provides an interface, coined as the *user interface*. In the older days this was done by typing some cryptic commands into a typewriter of a console device. Time changed and the first computer with a *Graphical User Interface (GUI)* emerged. A GUI is a visual interface to the user where the screen can host several windows each dedicated to a different tasks. Elements of the operating system (processes, files, directories, devices, network communicates) and their status are symbolized by icons and the interactions is mostly by moving and clicking a pointing device which is another icon (usually a movable arrow) on the screen. The [Xerox Alto](http://www.xerox.com)⁵ introduced on March 1973, was the first computer that had a GUI. The first personal computer with a GUI was [Apple Lisa](http://www.apple.com)⁶, introduced in 1983 with a price of US\$10000. Almost three years later, by the end of 1985 Microsoft released its first OS with GUI: Windows 1.0. The archaic console typing still exists, in the form a type-able window, which is still very much favored among programming professionals.

Exercise

- To gain more insight, play around with the Von Neuman machine simulator at <http://vnsimulator.altervista.org>
- Using Google and the manufacturer's web site, find the following information for your desktop/laptop:
 - Memory (RAM) size
 - CPU type and Clock frequency
 - Data bus size
 - Address bus size
 - Size of the general purpose registers of the CPU
 - Harddisk or SSD size and random access time

⁵ <https://github.com/sinankalkan/CENG240/blob/master/figures/XeroxAlto.jpg?raw=true>

⁶ <https://github.com/sinankalkan/CENG240/blob/master/figures/AppleLisa.jpg?raw=true>

Opcode	Affected address
0001	0100

- Assume that we have a CPU that can execute instructions with the following format and size given above.
 - What is the number of different instructions that this CPU can decode?
 - What is the maximum number of rows in the memory that can be addressed by this CPU?

Further Reading

- Computer Architectures:
 - Von Neumann Architecture: http://en.wikipedia.org/wiki/Von_Neumann_architecture
 - Harvard Architecture: http://en.wikipedia.org/wiki/Harvard_architecture
 - Harvard vs. Von Neumann Architecture: http://www.pic24micro.com/harvard_vs_von_neumann.html
 - Quantum Computer: http://en.wikipedia.org/wiki/Quantum_computer
 - Chemical Computer: http://en.wikipedia.org/wiki/Chemical_computer
 - Non-Uniform Memory Access Computer: http://en.wikipedia.org/wiki/Non-Uniform_Memory_Access